

Founder's Office (Marketing & AI)

Take-Home Assignment

Build a working AI agent that, every week:

1. Finds the best-performing content on X in our industry
2. Analyzes why it performed
3. Generates content ideas we can actually post

Build it in Claude Code.

The Context:

- We're a revenue-focused UX design agency.
- Our clients: Venture-backed SaaS and AI companies (seed to Series A stage)
- Our growth channel: X conversations with SaaS founders, product/design people, build-in-public folks, AI builders
- Our content focuses on insights, lessons, case studies, and design thinking that resonate with this audience.

What the agent should do?

On a weekly cadence, the agent should:

STEP 1: RESEARCH

Pull the top-performing posts from your chosen set of accounts and topics over the past week.

You define "our industry" on X. Pick a set of accounts and topics you think are the right ones to learn from. Document your logic. (Choosing the right inputs is part of the work.)

STEP 2: ANALYZE

Rank the top posts against a sensible benchmark.

(Think: engagement relative to profile size, not raw likes)

Figure out WHY the best ones worked:

- The hook

- The format
- The angle
- The structure
- The timing
- The pattern

STEP 3: IDEATE

Produce content ideas for us to post.

For each idea:

- Point to the post or pattern it's modeled on
- Explain why that pattern worked
- Adapt it to our voice (not a copy)
- Make it distinctive to Bricx Labs

STEP 4: PACKAGE

Hand back a weekly brief that a content writer and Twitter strategist can execute without needing to ask for clarification.

Output should be:

- Clear and actionable
- Self-contained
- Ready to publish (with minor edits)

GETTING X DATA

You'll need a way to access X data. Options:

- Official X API
- Third-party tools (e.g., Apify)
- Other approaches you find

Working out a usable method is part of the exercise. We've left it open on purpose.

IMPORTANT: We care far more about the system you build than about perfect live data.

If you hit a wall on API access:

- ✓ A working pipeline on sample data you've assembled is completely fine
- ✓ Just tell us how you'd productionize it

- ✓ Bonus points for working around constraints creatively

WHAT TO SUBMIT?

Submit THREE things:

1. THE AGENT

- Code, prompts, workflow (however you've structured it)
- Quick notes on how to run it
- How to use/execute weekly
- Tools and dependencies

2. A VIDEO WALKTHROUGH (3-5 minutes, Loom)

- Walk us through what you built
- Explain where a human stays in the loop
- Tell us what breaks or needs supervision
- Share your thinking and approach
- (This is important—we want to hear your reasoning)

3. ONE SAMPLE OUTPUT

- One complete weekly output showing:
 - * The research/analysis
 - * The content ideas (with patterns and reasoning)
 - * The executable brief
- Show us what the agent actually produces

WHAT GOOD LOOKS LIKE

- ✓ It runs. The weekly output makes sense.
- ✓ Analysis explains WHY top content performed, not just which posts got engagement.
- ✓ Each idea traces back to a pattern and a reason.
- ✓ Ideas read as on-brand, not generic.
- ✓ The final brief is genuinely executable. A content writer could run with it cold.
- ✓ You've thought about the system, not just the output.

WHAT WE'RE EVALUATING

Execution and shipping (40%)

Thinking and judgement (35%)

System design (15%)
Communication (10%)

All the best!